

PlanTrace-VL: auditable vision-language grounding for document-centered planning review

Yuhang Yan¹, Xuning Tan¹, and Zhenkai Gao^{*2}

¹School of Science and Engineering, The Chinese University of Hong Kong, Shenzhen, Shenzhen, China

²School of Advanced Technology, Xi'an Jiaotong-Liverpool University, Suzhou, China

ABSTRACT

Vision-language and language models are increasingly used to support urban planning tasks such as plan reading, map interpretation, and planning document analysis. Yet professional planning review is not a single-image visual question-answering task. It is document-centered, evidence-intensive, and accountable. Planners inspect review packets that combine zoning maps, site plans, permit drawings, legends, annotations, dimensional markings, frontage information, parcel context, and regulatory excerpts. Their findings must remain traceable to evidence and subject to professional judgment. Direct answer generation is therefore insufficient on its own: a useful assistant should expose the visual evidence used, the planning facts extracted, the regulatory clauses linked, and the issues that require human verification. We present PlanTrace-VL, a document-centered, human-in-the-loop framework for auditable planning document review. Rather than treating vision-language models as autonomous approval systems, PlanTrace-VL positions them as evidence-grounding assistants for document-centered planning review. The framework supports evidence localization, planning fact extraction, regulatory clause linking, review-packet generation, and human-verification status prediction, with outputs designed to distinguish supported findings from cases requiring planner verification.

To train and evaluate the framework, we construct UrbanPlanTrace-Review, a curated open-source proxy benchmark and task suite for auditable planning document review. The benchmark organizes visual grounding, regulatory clause linking, review-packet generation, and cross-page consistency tasks around traceable intermediate outputs. PlanTrace-VL is trained with a two-stage LoRA-based supervised fine-tuning strategy: visual adaptation for technical drawing understanding followed by regulatory-clause-grounded instruction tuning for regulatory-clause-linked review tasks. In a 300-sample-per-task evaluation, PlanTrace-VL reaches 95.67 Evidence Localization and 92.89 Fact F1, while exposing remaining weaknesses in regulatory clause linking, unsupported findings, and cross-page consistency. These findings support a bounded contribution: specialized vision-language training can improve auditable intermediate outputs for planning review while preserving human planner verification and judgment.

Keywords: Vision-language models, Urban planning document review, Human-AI collaboration, Urban planning, Multimodal evidence grounding

INTRODUCTION

Computer-based planning tools have long been framed as systems that support, rather than replace, professional reasoning (Klosterman, 1997; Geertman and Stillwell, 2004). Recent language-model research extends this tradition to plan reading, policy analysis, public participation, and planning decision support (Peng et al., 2024; Fu, 2024; Fu et al., 2024). Professional planning review, however, poses a particularly demanding support problem. The unit of analysis is often a multimodal case packet containing zoning maps, site plans, permit drawings, legends, annotations, dimensions, frontage and parcel information, and regulatory excerpts. A defensible finding must therefore remain connected to the document evidence and regulatory provision on which it rests.

Visually rich document understanding provides useful foundations for this setting. DocVQA established question answering over document images, while DUDE and GRAM extended evaluation and modeling toward diverse, multi-page documents (Mathew et al., 2021; Van Landeghem et al., 2023; Blau et al., 2024). Domain-specific vision-language models have also advanced planning-map interpretation




Figure 1. Conceptual motivation of PlanTrace-VL.

(Zhu et al., 2025). Yet answer-oriented document and map benchmarks do not by themselves represent the evidentiary workflow of planning review: locating the relevant region, extracting a planning fact, linking that fact to a candidate clause, recording whether the link is supported, and flagging the result for human verification. Automated compliance-checking research similarly shows that regulatory review requires explicit representations of both design information and rules (Solihin and Eastman, 2015; Zhang and El-Gohary, 2016, 2017). The open problem is thus not simply multi-page question answering, but auditable grounding across heterogeneous planning evidence and regulatory text.

Reliability makes this distinction consequential. Large vision-language models can generate image-inconsistent objects and conclusions, and dedicated benchmarks continue to expose visual hallucination and image-context reasoning failures (Li et al., 2023; Guan et al., 2024). A fluent review note can therefore appear resolved even when its evidence or regulatory link is weak. Human-AI interaction research recommends interfaces that preserve user control, communicate uncertainty, and discourage uncritical reliance (Amershi et al., 2019; Bućinca et al., 2021; Mosqueira-Rey et al., 2023). For planning review, these principles imply a concrete technical requirement: the model should expose inspectable intermediate artifacts instead of hiding them behind a final answer. Figures 1 and 2 summarize this shift from answer generation to document-centered, auditable review.

To meet this requirement, we introduce *PlanTrace-VL*, a human-in-the-loop framework for auditable vision-language grounding in planning document review. PlanTrace-VL treats the model as a review-support system that exposes evidence rather than makes final decisions. It localizes visual evidence in maps and drawings, extracts planning facts such as land-use labels and dimensional notes, links those facts to candidate regulatory clauses, and drafts review notes that separate supported findings from cases requiring human verification. Its scope is deliberately bounded to review support; automated approval, complete legal compliance checking, and professional judgment remain outside the system’s claims.

We also introduce UrbanPlanTrace-Review, a curated open-source proxy benchmark and task suite for training and evaluating this workflow. The benchmark covers visual grounding, regulatory clause linking, review-packet generation, and cross-page consistency. PlanTrace-VL uses a two-stage LoRA-based supervised fine-tuning strategy: technical-drawing adaptation first, followed by regulatory-clause-grounded instruction tuning. At inference time, layout-aware patch extraction, planning fact extraction, hierarchical clause retrieval, clause-link consistency checks, and structured review-note generation produce a trace that a planner can inspect.




Figure 2. UrbanPlanTrace-Review benchmark showcase for auditable planning document review. The left panel shows traceable review tasks built from local prediction records and public drawing resources. The middle panel summarizes non-exclusive diagnostic error rates from the PlanTrace-VL evaluation logs. The right panel compares auditable review scores from logged evaluation metrics. The figure illustrates the paper’s evaluation logic: models are assessed through evidence, regulatory clause links, review traces, diagnostic risks, and auditable scores rather than autonomous approval decisions.

The evaluation measures each layer of this trace rather than only final-answer accuracy. On four 300-sample diagnostic tasks, PlanTrace-VL reaches 95.67% Evidence Localization and 92.89% Fact F1, the strongest results in the reported comparison. The remaining results also delimit the contribution: Consistency F1 is only 19.66%, the unsupported-link rate is 19.67%, and the model does not reliably distinguish insufficient visual evidence from other verification cases. These outcomes support a focused claim—domain adaptation improves several traceable intermediate outputs—without establishing municipal deployment readiness or autonomous review capability.

This paper makes three contributions. First, it formulates planning document review as an auditable evidence-grounding workflow rather than a single-answer prediction task. Second, it introduces UrbanPlanTrace-Review and a two-stage training pipeline that connects technical-drawing understanding with regulatory-clause-grounded review. Third, it provides a layer-wise diagnostic protocol that identifies both improvements in traceable intermediate outputs and failure modes that still require planner oversight.

RELATED WORK

LLMs and AI-Assisted Urban Planning

Planning support systems traditionally combine data, models, and visualization to assist professional planning activities (Klosterman, 1997; Geertman and Stillwell, 2004). Language models extend this agenda to plan reading, policy synthesis, public participation, and decision assistance, while retaining concerns about professional judgment, public accountability, and institutional context (Peng et al., 2024). NLP studies analyze resilience plans, general plans, and policy documents at scale to reveal themes, omissions, and jurisdictional variation (Fu et al., 2023; Brinkley and Stahmer, 2024; Fu, 2024). More recent work tests whether general-purpose LLMs can evaluate plans or support participatory workflows (Fu et al., 2024; Zhou et al., 2024b). These studies establish the value of language-based planning assistance, but they primarily analyze text or planning knowledge. PlanTrace-VL instead targets review packets in which visual evidence and regulatory text must be connected through an inspectable trace.

Urban Planning Benchmarks and Urban Intelligence Agents

Urban planning and urban computing benchmarks test whether foundation models can reason about cities as social, spatial, and temporal systems. UrbanPlanBench evaluates planning principles, professional knowledge, and regulations (Zheng et al., 2025); CityBench and USTBench address broader urban reasoning, world modeling, and spatio-temporal prediction (Feng et al., 2024; Lai et al., 2025). Agent-

oriented platforms additionally study how language-model agents act in or reason about simulated city environments (Xu et al., 2023). These benchmarks move beyond generic NLP, but their principal targets remain planning knowledge, urban reasoning, prediction, or agent behavior. UrbanPlanTrace-Review evaluates a different workflow: finding document evidence, extracting review-relevant facts, linking them to candidate clauses, and preserving a trace for planner verification.

Visually Rich and Multi-Page Document Understanding

Document AI provides the closest general formulation of the paper’s multimodal input problem. DocVQA established visual question answering over document images and highlighted the importance of document structure (Mathew et al., 2021). DUDE broadened this setting to diverse, multi-domain, and multi-page documents, while GRAM introduced document-level reasoning across pages (Van Landeghem et al., 2023; Blau et al., 2024). This literature demonstrates that text, layout, and visual content must be modeled jointly and that page-level processing is insufficient for many real documents. Its typical output, however, remains an answer to a question. PlanTrace-VL adopts multi-page document reasoning as a foundation but evaluates a planning-specific chain of localized evidence, extracted facts, candidate regulatory clauses, support status, and review notes.

Vision-Language Models for Maps and Urban Visual Materials

Vision-language research provides multimodal foundations for interpreting maps, remote-sensing imagery, urban scenes, and planning materials. RSVQA, GeoChat, RemoteCLIP, EarthGPT, and GEOBench-VLM evaluate geospatial question answering, image–text alignment, localization, and spatial interpretation (Lobry et al., 2020; Kuckreja et al., 2023; Liu et al., 2023; Zhang et al., 2024; Danish et al., 2025). UrBench extends multimodal evaluation to multi-view urban scenarios (Zhou et al., 2024a). Most directly, PlanGPT-VL adapts vision-language modeling to planning-map interpretation and uses structured verification to reduce hallucination (Zhu et al., 2025). PlanTrace-VL builds on this domain-specific direction but changes the unit of evaluation from a map answer to an evidence chain spanning documents, facts, candidate clauses, and review status.

Technical Drawing and Planning-Document Understanding

Planning drawings differ from natural images because symbols, linework, scales, labels, legends, and dimensions jointly encode design meaning. FloorPlanCAD and ArchCAD-400K provide large-scale resources for panoptic symbol spotting in CAD drawings, while CubiCasa5K supports recognition of rooms, walls, openings, and related floor-plan elements (Fan et al., 2021; Luo et al., 2025; Kalervo et al., 2019). Specialized symbol-spotting research likewise shows that architectural semantics require drawing-aware visual grounding (Rezvanifar et al., 2020). TechMB extends this concern to vision-language interpretation of technical drawings (Kunz et al., 2025). These resources motivate PlanTrace-VL’s first-stage visual adaptation. They do not address the subsequent planning-review problem of converting localized evidence into facts, regulatory links, and verification-aware notes.

Regulatory Clause Grounding and Automated Compliance Checking

Planning document review is also a regulatory interpretation task. Land-use labels, frontage conditions, setbacks, height limits, lot coverage, overlays, and parcel boundaries matter because zoning and development codes translate planning goals into review criteria (Talen, 2012). Automated compliance checking represents rules and design information computationally through rule classification, natural-language information extraction, logic-based reasoning, BIM-based checking, and machine-readable workflows (Dimyadi and Amor, 2013; Solihin and Eastman, 2015; Zhang and El-Gohary, 2016, 2017; Demir Altintas and Ilal, 2022). CODE-ACCORD, the National Zoning and Land Use Database, and the National Zoning Atlas further illustrate the growing availability of structured or semi-structured regulatory data (Hettiarachchi et al., 2025; Mleczo and Desmond, 2023; National Zoning Atlas, 2026). PlanTrace-VL uses these resources to surface candidate provisions and expose the evidence behind a link. It does not claim autonomous planning approval or definitive legal compliance checking.

Auditable and Human-in-the-Loop Planning AI

Because planning review affects development rights and public accountability, AI assistance must be reviewable rather than merely fluent. Human–AI interaction research emphasizes user control, uncertainty communication, and appropriate reliance (Amershi et al., 2019; Mosqueira-Rey et al., 2023). Reviewable




Figure 3. Overview of the PlanTrace-VL framework.

automated decision-making similarly calls for records that can be inspected, contested, and governed (Cobbe et al., 2021), while evaluative AI emphasizes evidence for and against a decision rather than a one-shot recommendation (Miller, 2023; Le et al., 2024). This need is reinforced by documented hallucination and visual-reasoning failures in large vision-language models (Li et al., 2023; Guan et al., 2024) and by evidence that uncritical reliance can persist in AI-assisted decisions (Buçinca et al., 2021). PlanTrace-VL operationalizes these principles through evidence localization, fact extraction, clause relevance, unsupported-link detection, review-note factuality, human-verification behavior, and cross-page consistency.

Together, prior work advances planning LLMs, urban benchmarks, multi-page document understanding, map-oriented VLMs, drawing interpretation, regulatory rule representation, and human-centered AI. PlanTrace-VL connects these strands around a narrower problem: producing an auditable planning-review trace from document evidence to planner-verifiable notes. The methodology below translates this synthesis into a proxy benchmark, training strategy, inference pipeline, and diagnostic evaluation.

METHODOLOGY

Design Principle: Document-Centered Auditable Review

PlanTrace-VL treats planning document review as a review-packet task rather than a single-image visual question-answering task. In practice, a planner may inspect zoning maps, site plans, permit drawings, legends, annotations, dimensional notes, frontage information, parcel boundaries, and regulatory excerpts as one evidentiary record. The methodological objective is to generate traceable intermediate artifacts that a planner can inspect, compare, and revise, not to produce an autonomous approval decision. This design principle determines the rest of the pipeline: the benchmark must contain review-relevant evidence, the tasks must expose intermediate outputs, and the inference schema must preserve uncertainty. PlanTrace-VL is therefore organized around auditable grounding, where each review note should be traceable to localized visual evidence, extracted planning facts, and candidate regulatory clauses.

UrbanPlanTrace-Review Construction

Real municipal permit-review packets are often sensitive, inconsistent, and not publicly redistributable. We therefore construct UrbanPlanTrace-Review as a curated open-source proxy benchmark and task suite for auditable planning document review. The benchmark combines public architectural drawings, technical diagrams, floor-plan resources, and structured regulatory text to approximate review-relevant visual and textual signals. This construction supports reproducibility and allows the data and task design to be inspected. It also defines a central boundary of the study: UrbanPlanTrace-Review is a proxy benchmark for studying auditable planning review, not a complete representation of municipal permit-review practice, local administrative discretion, or official legal review records. Because its visual resources include substantial technical drawing and CAD-style material, the benchmark likely overrepresents drawing-element localization relative to full urban planning review packets. It should therefore be interpreted as a controlled diagnostic setting for traceable review behavior, not as evidence of municipal deployment readiness.

Figure 4 summarizes UrbanPlanTrace-Review at the corpus level. The distribution separates visual adaptation, regulatory-clause-grounded review tuning, and diagnostic review tasks. The accompanying statistics report global corpus properties that complement, rather than duplicate, the task-split table.




Figure 4. UrbanPlanTrace-Review dataset overview. The ring chart shows the distribution of curated task instances used to construct the benchmark, and the table reports global corpus statistics that complement the detailed train/evaluation split in Table 1. Percentages in the ring chart are computed over task instances; zoning/code records are reported separately because they form the retrieval index rather than supervised task examples.

The corpus-level summary is complemented by sample-level examples in Figure 5. These examples clarify how UrbanPlanTrace-Review differs from a conventional visual question-answering benchmark. Each record is organized around an auditable review step: a document input, a review-oriented question, a gold structured output, and a trace note that preserves the need for planner verification.




Figure 5. Representative UrbanPlanTrace-Review samples across auditable planning-review tasks. The examples illustrate evidence localization, planning fact extraction, regulatory clause linking, review-packet generation, cross-page consistency checking, and human-verification status prediction. Each panel is drawn from logged task records using image samples that do not overlap with Figure 2. Green text marks gold structured outputs, not final planning approvals.

Visual Grounding Resources

Visual grounding in UrbanPlanTrace-Review draws on public drawing resources, including ArchCAD-400k, FloorPlanCAD, TechMB, and related floor-plan repositories. These resources support technical drawing adaptation by exposing the model to dense visual materials with symbols, legends, annotations, structural elements, room or object labels, dimensional markings, and layout patches. In planning review, such elements function as candidate evidence for planning facts, such as labeled uses, frontage conditions, setback-related markings, dimension notes, or regions requiring human checking. The visual grounding component is designed to identify where review-relevant evidence appears and convert it into structured facts for downstream inspection.

Rule-Grounding Resources

Rule grounding in UrbanPlanTrace-Review incorporates public regulatory resources, including CODE-ACCORD, the National Zoning Atlas, NZLUD, and structured zoning text. These resources provide regulatory clauses, zoning-related text, and structured land-use information for retrieval and regulatory clause linking tasks. From a zoning-and-codes perspective, land-use labels, frontage conditions, setbacks, dimensional controls, overlays, height or coverage limits, and parcel-specific requirements are review-relevant because they connect visual or textual facts to enforceable standards. PlanTrace-VL treats this connection as regulatory-clause-linked evidence support. A retrieved provision is not treated as a final legal determination, and UrbanPlanTrace-Review is not assumed to capture all local exceptions, variances, administrative interpretations, or discretionary procedures.

Task Construction

The UrbanPlanTrace-Review tasks operationalize the design principle by converting planning review into a sequence of inspectable subtasks. Table 1 reports the train, validation, and evaluation counts used in the reported experiments. Each task has a defined input, expected output, and review purpose, allowing performance to be interpreted as evidence about a specific layer of audibility.

Evidence localization. The input is a map, drawing, diagram, or layout patch. The expected output is a localized evidence description, region label, or patch-level grounding target. The review purpose is to

indicate where a planner should look before accepting an extracted fact or review note.

Planning fact extraction. The input is localized visual evidence from a drawing or map region. The expected output is a structured planning fact, such as a land-use label, frontage condition, dimensional note, setback-related marking, or relevant annotation. The review purpose is to convert visual evidence into a factual statement that can be checked by a reviewer.

Regulatory clause linking. The input is an extracted planning fact and a set of candidate regulatory clauses. The expected output is a linked clause or ranked set of relevant clauses. The review purpose is to connect document evidence to regulatory text that may support or constrain the review finding.

Review-packet generation. The input is a proxy packet containing visual evidence, extracted facts, and candidate clauses. The expected output is a structured review note that separates supported findings from uncertain or unsupported statements. The review purpose is to organize review artifacts, not to produce an unverified final decision.

Cross-page consistency checking. The input is a multi-page proxy packet or a set of related evidence items across documents. The expected output is a consistency status or inconsistency flag. The review purpose is to detect whether planning facts, labels, notes, and regulatory clause links agree across pages.

Human-verification status prediction. The input is a generated review trace containing evidence, facts, clauses, and a draft note. The expected output is a human-verification status, such as supported, uncertain, or requiring planner review. The review purpose is to prevent weakly supported outputs from being treated as authoritative.

Two-Stage LoRA-Based Supervised Fine-Tuning

Training follows a two-stage LoRA-based supervised fine-tuning strategy. Stage 1 performs visual adaptation for technical drawing understanding. In this stage, the model is tuned to parse floor plans, drawing symbols, legends, annotations, structural elements, layout patches, and dimensional markings. This stage addresses the visual conventions of planning-related drawings, which differ from natural images and generic VQA benchmarks.

Stage 2 performs regulatory-clause-grounded instruction tuning for regulatory-clause-linked review tasks. In this stage, the model is trained on examples that combine localized evidence, extracted planning facts, candidate clauses, review-note structures, and human-verification labels. The stage order is deliberate: the model first learns to identify technical visual evidence, then learns to connect that evidence to regulatory clauses and review notes. This sequence is intended to reduce the risk that outputs involving regulatory clauses are generated without stable visual support.

Inference Pipeline

Figure 3 summarizes the inference workflow. First, layout-aware patch extraction identifies document regions likely to contain review-relevant evidence. Second, evidence localization assigns candidate evidence descriptions or region labels to those patches. Third, planning fact extraction converts localized evidence into structured planning facts. Fourth, hierarchical clause retrieval searches the regulatory text index to surface candidate provisions at the corpus, document, and clause level. Fifth, consistency checking for regulatory clause linking compares extracted facts with retrieved clauses and flags weak or mismatched links. Sixth, review-note generation drafts a structured note that records localized evidence, extracted facts, linked regulatory clauses, and verification status. Finally, a human planner reviews the trace, accepts or revises supported findings, and escalates uncertain or unsupported items. The model generates candidate evidence chains and notes; interpretation, escalation, and professional judgment remain under planner oversight.

Output Schema and Auditability

PlanTrace-VL defines auditability through the structure of its outputs. Each generated review trace is expected to contain localized evidence, extracted facts, linked regulatory clauses, supported findings, unsupported or uncertain findings, human-verification flags, and review notes. These fields correspond to review actions in professional planning practice: locating evidence, checking factual claims, comparing a claim with applicable standards, identifying unsupported links, and deciding whether a case requires further human review. This schema is more inspectable than a single final answer because it allows

planners to follow the reasoning path, identify where an error may have entered, and preserve an accountable review record. It also determines how the experiments are evaluated, since each metric targets one layer of the trace rather than a single end label.

Claim Boundary

This study uses UrbanPlanTrace-Review to examine auditable evidence grounding in planning document review. UrbanPlanTrace-Review is a curated open-source proxy benchmark and task suite rather than a municipal permit-review corpus. PlanTrace-VL does not claim to automate official planning approval, replace professional judgment, or fully replicate municipal review practice. It should be understood as a research framework that generates traceable intermediate outputs for planner verification.

EXPERIMENTS

Diagnostic Evaluation Design

The experiments evaluate PlanTrace-VL as a diagnostic system for auditable planning document review. The evaluation target is review support, not autonomous approval or rejection of planning documents. It follows the review sequence that a planner would need to inspect: localizing visual evidence, extracting planning facts, linking those facts to regulatory clauses, generating structured review notes, preserving human-verification flags, and checking cross-page consistency. Each experiment therefore corresponds to a review-support step. Visual grounding tests whether evidence can be traced to a document region. Planning fact extraction tests whether drawings and annotations can be converted into reviewable statements. Regulatory clause linking tests whether those statements can be connected to candidate regulatory provisions. Review-packet generation tests whether the output remains parseable, factual, and appropriately cautious. The results are interpreted as evidence about auditable review support, not as a generic model leaderboard.

Evaluation Tasks and Data Splits

Table 1 reports the UrbanPlanTrace-Review task construction and data splits used in the reported PlanTrace-VL experiments. Each split corresponds to a review action in document-centered planning practice. Stage 1 visual SFT supports technical drawing adaptation, so the model can interpret dense visual signals such as symbols, layout patches, legends, annotations, and dimensional markings. Stage 2 rule/review SFT supports regulatory-clause-grounded review tuning, where extracted planning facts must be connected to regulatory text and expressed as inspectable review notes. The rule corpus index provides candidate zoning and code provisions for retrieval. The retrieval rule split tests whether the relevant provision can be surfaced before generation. The layout patch grounding split evaluates whether visual evidence can be localized and converted into planning facts. The review packet proxy split evaluates structured review-note generation. The cross-page consistency split tests whether distributed evidence can be reconciled across a proxy packet.

Table 1. UrbanPlanTrace-Review task construction used in the reported PlanTrace-VL experiments. Counts are taken from the local UrbanPlanTrace-Review manifests and training configurations.

Task/data split	Main source signal	Train / pool	Eval / val.	Split source	Supported evaluation
Stage 1 visual SFT	ArchCAD / FloorPlanCAD / TechMB	26,270 pool	789 val.	runtime split	Technical drawing adaptation
Stage 2 rule/review SFT	Structured zoning text + review packets	6,021 pool	302 val.	runtime split	Regulatory-clause-grounded review tuning
Rule corpus index	Zoning and code text	9,677 docs	N/A	retrieval index	Candidate rule retrieval
Retrieval rule	Rule candidates + gold clause	1,200 train	300 eval	manifest split	Clause retrieval
Layout patch grounding	Drawing/map patch + localized labels	1,200 train	300 eval	manifest split	Evidence localization, planning fact extraction
Review packet proxy	Image/facts + candidate clauses	1,200 train	300 eval	manifest split	Review packet quality
Cross-page consistency	Multi-page proxy packet	1,200 train	300 eval	manifest split	Cross-page consistency

Baselines and Fairness

PlanTrace-VL is compared against prompting-based VLM/LLM baselines reported in Table 2. The evaluated baseline suite covers Mistral-Small-3.2-24B-Instruct-2506 (Mistral AI, 2025; Mistral AI Team, 2025), MiniCPM-V-2.6-8B (Yao et al., 2024), Gemma-3-27B/12B/4B-IT (Gemma Team et al., 2025), Qwen3-VL-32B/30B-A3B/8B/4B-Instruct (Bai et al., 2025b), and Qwen2.5-VL-72B/32B/7B/3B-Instruct (Bai et al., 2025a). The comparison is diagnostic rather than leaderboard-oriented. It tests whether domain

adaptation and structured outputs improve auditability for planning review, not whether the largest model has the strongest generic visual question-answering ability. This distinction matters because a model may answer an isolated clause-selection question correctly while failing to expose localized evidence, trace fields, uncertainty status, or review notes that a planner can verify. For transparency, the baseline run produced all 52 expected prediction files with 300 rows per file and no malformed JSONL rows; the aggregate is still marked incomplete because it contains 15 sample-level API errors. The PlanTrace-VL evaluation is complete with 0 errors across the four scored diagnostic tasks.

Metrics as Auditability Layers

We group the metrics into auditability layers rather than treating them as a single accuracy score. Some metrics are reported as headline component scores, while others appear in diagnostic tables or as source metrics for error analysis. The evidence layer includes Evidence Localization and Evidence Traceability, which indicate whether a planner can find the visual source behind a generated claim. The fact layer uses Fact F1 to evaluate whether review-relevant facts are extracted from drawings, maps, annotations, and dimensional notes. The rule layer includes Rule Link Accuracy, Clause Relevance, and Unsupported-Link Rate, because zoning and development review require facts to be connected to applicable provisions without unsupported evidence or unsupported clause links. The review-note layer includes JSON Validity, Note Factuality, and False-Supported Rate, the packet-level risk that a note presents an unsupported finding as supported. These metrics assess whether the generated packet is parseable, factually grounded, and cautious about unsupported findings. The oversight layer uses Human Verification to test whether the system preserves an escalation path for planner review. The cross-page layer uses Consistency F1 and cross-page status calibration, reflecting the fact that planning evidence is often distributed across more than one page, drawing, or regulatory excerpt.

Component-Level Results

Table 2 reports component-level performance on the 300-sample-per-task evaluation. PlanTrace-VL shows its clearest gains on structured intermediate outputs that are central to the paper’s claim. It reaches 95.67 Evidence Localization and 92.89 Fact F1, outperforming the prompting baselines on these two review-support functions. It also obtains the highest Consistency F1 in the table at 19.66, although the absolute value remains low. This result indicates that cross-page reasoning remains difficult in the proxy setting. Regulatory clause linking results are more mixed. PlanTrace-VL reaches 95.67 Rule Link Accuracy, but larger prompting baselines such as Qwen2.5-VL-72B-Instruct and Qwen2.5-VL-32B-Instruct remain competitive or slightly higher on isolated Rule Link Accuracy. Unsupported-Link Rate also requires cautious interpretation. PlanTrace-VL reduces this rate relative to several strong baselines with high unsupported-link rates, but it is not the lowest value in the table. These patterns support a bounded conclusion: the main benefit of PlanTrace-VL is improved traceable intermediate work, not universal dominance on every isolated metric.

Table 2. Component-level performance on the reported 300-sample-per-task evaluation. Values are percentages. Higher is better except unsupported-link rate.

Model	Evidence Loc.	Fact F1	Rule Link Acc.	Consistency F1	Unsupported-Link Rate ↓
Mistral-Small-3.2-24B-Instruct-2506	73.33	62.58	93.33	2.14	24.00
MiniCPM-V-2.6-8B	64.67	63.27	19.00	0.42	39.67
Gemma-3-27B-IT	59.67	56.08	90.67	1.33	0.00
Gemma-3-12B-IT	43.00	55.05	91.33	1.21	0.00
Gemma-3-4B-IT	37.67	10.48	41.00	2.75	0.33
Qwen3-VL-32B-Instruct	42.67	37.42	0.00	1.40	0.00
Qwen3-VL-30B-A3B-Instruct	53.00	48.58	5.00	2.67	0.00
Qwen3-VL-8B-Instruct	62.00	50.95	0.00	1.06	0.00
Qwen3-VL-4B-Instruct	1.00	0.73	0.00	0.25	0.00
Qwen2.5-VL-32B-Instruct	58.67	73.33	96.33	1.28	39.00
Qwen2.5-VL-72B-Instruct	81.67	58.18	96.67	8.56	39.67
Qwen2.5-VL-7B-Instruct	93.00	61.98	90.67	11.89	39.67
Qwen2.5-VL-3B-Instruct	66.33	81.02	69.00	8.42	39.33
PlanTrace-VL	95.67	92.89	95.67	19.66	19.67

End-to-End Review Packet Quality

The final column of Table 2 evaluates a central risk in end-to-end proxy review packets: unsupported links between a generated finding and its supposed evidence or clause support. This metric is more review-facing than JSON validity alone. A fluent review note can be well formatted but still unreliable for review if it incorrectly states support. PlanTrace-VL reduces this unsupported-link rate to 19.67, below several strong prompting baselines, although the result should be read diagnostically rather than as evidence of autonomous review readiness. Low values for models with poor packet parseability should not be interpreted as evidence of reliable review behavior.

Diagnostic Error Analysis

Table 3 translates PlanTrace-VL’s remaining metric failures into review-facing risks. Unlocalized or mismatched evidence means that a note may lack visual support. Wrong or missing regulatory clause links mean that an applicable zoning provision may be misidentified. False-supported findings are especially important in planning review because they can make an uncertain case appear resolved. Invalid proxy JSON prevents downstream inspection of trace fields. Cross-page finding mismatch shows that distributed evidence remains difficult to reconcile. These errors should be read as audit signals, not only as aggregate performance losses. They identify where the system still requires human oversight and where future work should strengthen the pipeline.

Table 3. Diagnostic error analysis for the PlanTrace-VL run. Values are derived from the reported metric outputs and shown as percentages.

Diagnostic error	Source metric	Rate	Review implication
Unlocalized or mismatched evidence	100 - proxy evidence F1	44.67	Review note lacks traceable visual support
Wrong or missing clause link	100 - proxy clause accuracy	52.00	Applicable regulatory clause may be misidentified
False-supported finding	Proxy False-Supported Rate	19.67	Output states support without evidence
Missing human-verification flag	100 - Human Verification rate	0.00	Planner oversight flag is preserved in this run
Invalid proxy JSON	100 - proxy JSON validity	1.00	Trace fields cannot be parsed reliably
Cross-page finding mismatch	100 - cross-page finding F1	80.34	Distributed evidence is not reconciled

Evidence-Type Difficulty

Table 4 stratifies the layout-patch grounding task by visual evidence type. Aggregate scores can hide whether performance is stable across the categories that planners inspect in maps and drawings. In the current proxy evaluation, many structural drawing elements receive high localization scores, including Wall, Axis & Grid, Concrete Beam, Foundation, Glass, Steel Beam, Parking Space, Rebar, and Sink. Smaller or more visually ambiguous elements, such as Hole and Pile, show lower evidence scores. These evidence types should be interpreted as proxy visual categories for testing localization and traceability, not as a complete representation of urban planning judgment. This matters for planning review because strong performance on large structural elements does not guarantee reliable grounding of small symbols, parcel annotations, overlay boundaries, frontage labels, or setback-related markings.

Table 4. Evidence-type difficulty on the layout-patch grounding task. Values are percentages except n and are computed from the reported PlanTrace-VL prediction file.

Evidence type	n	Target Acc.	Evidence F1	Fact F1
Concrete Column	50	92.00	92.00	90.68
Wall	49	100.00	100.00	98.75
Axis & Grid	36	100.00	100.00	88.00
Concrete Beam	36	100.00	100.00	93.65
Foundation	13	100.00	100.00	90.97
Hole	13	84.62	84.62	89.35
Pile	13	100.00	84.62	92.64
Glass	11	100.00	100.00	98.60
Steel Beam	10	100.00	100.00	95.14
Parking Space	9	100.00	100.00	100.00
Rebar	8	100.00	100.00	94.23
Sink	8	100.00	100.00	97.50

Review-Status Calibration

Table 5 examines whether PlanTrace-VL preserves the human-verification pathway, a central safeguard in planning review. A cautious assistant that escalates uncertain cases is preferable to one that falsely states support without evidence. The model achieves high Status Accuracy for “Requires human verification”, reaching 98.90, but performs poorly on “Insufficient visual evidence”, where Status Accuracy is 0.00 and False-Supported Rate is 49.58. It therefore preserves escalation more reliably than it distinguishes insufficient evidence from other verification cases. This behavior reinforces the need for human-in-the-loop use and should not be interpreted as autonomous approval capability.

Table 5. Proxy review-status error analysis for PlanTrace-VL. Values are percentages except n and are stratified by the gold review status.

Gold review status	n	Status Acc.	Evidence F1	Clause Acc.	False-Supported Rate ↓
Requires human verification	181	98.90	58.01	46.41	0.00
Insufficient visual evidence	119	0.00	51.26	50.42	49.58
Overall	300	59.67	55.33	48.00	19.67

Rule-Retrieval Ablation

Table 6 reports the rule-retrieval ablation used before response generation. Hierarchical retrieval improves the rank of the relevant planning provision relative to flat BM25: Recall@1 increases from 94.70 to 97.93, and MRR increases from 96.82 to 98.96. This improvement matters because cited clauses must be easy for planners to verify. In zoning and development review, a retrieved provision is a candidate clause rather than a final legal determination, and it should be inspected alongside the extracted fact and localized evidence.

Table 6. Rule-retrieval ablation on the hard regulatory clause linking evaluation set. Values are percentages except n and measure retrieval before response generation.

Retriever	n	Recall@1	Recall@3	Recall@8	MRR
Flat BM25	434	94.70	98.62	100.00	96.82
Hierarchical retrieval	434	97.93	100.00	100.00	98.96

Experiment-Level Diagnostic Visualizations

Figure 6 visualizes capability profiles and retrieval-rank behavior without repeating the raw table entries. The figure supports the diagnostic interpretation of the experiments: the framework is evaluated by its ability to support planner verification through localized evidence, organized facts, surfaced clauses, and escalation of uncertain cases. The visualization complements the tables by showing capability balance and retrieval behavior rather than adding another model ranking.




Figure 6. Experiment-level diagnostics for auditable planning document review. Left: capability profile across review functions, shown as within-metric percentiles. Right: gold-clause recall by retrieval-rank threshold, with an inset showing hierarchical filtering from the rule corpus to requirement-level candidates.

Evidence Status and Limitations

The experiments are based on UrbanPlanTrace-Review and its proxy review tasks. They support the claim that PlanTrace-VL improves several auditable intermediate outputs, especially evidence localization, planning fact extraction, and traceable packet generation. They do not establish readiness for municipal deployment, full legal compliance checking, or autonomous planning approval. Because the visual splits include substantial technical drawing and CAD-style evidence, the results should be interpreted as diagnostic evidence about auditable grounding rather than as a complete representation of municipal review practice. Future work should validate the framework with real planning review packets, jurisdiction-specific rules, planner-in-the-loop evaluation, and more diverse document types, including zoning maps, parcel-level diagrams, site plans, permit drawings, legends, annotations, frontage conditions, and dimensional controls.

DISCUSSION

A main implication of PlanTrace-VL is that multimodal AI for planning should be evaluated around review workflows, not only answer accuracy. Planning review is an accountability practice in which findings must be explained, checked, and justified. A planner may need to identify why a frontage condition matters, where a setback measurement appears, which clause is relevant, and why a case should be escalated rather than resolved. Traceability is therefore a first-order property of planning AI, not a presentation feature added after generation.

This framing clarifies the paper’s position within planning support research. PlanTrace-VL does not replace planning judgment with an algorithmic decision rule. It tests whether a vision-language model can make review artifacts more inspectable for planner verification. This distinction matters because planning decisions are shaped by local ordinances, discretionary procedures, public values, and institutional accountability. These conditions are not fully captured by a proxy benchmark, so the results should be interpreted as evidence about auditable review support rather than municipal deployment readiness.

The results are most meaningful where planning review requires structured intermediate outputs. PlanTrace-VL improves evidence localization and planning fact extraction, and it produces review packets with explicit trace fields. These gains are practice-facing because they help organize the materials that planners already need to inspect: localized evidence, extracted facts, linked regulatory clauses, and auditable notes. The value of PlanTrace-VL therefore lies in making the review record more inspectable, while leaving interpretation, discretion, and legal responsibility with human planners.

The remaining errors define important research boundaries. Unsupported findings, missing or wrong regulatory clause links, and cross-page mismatches are serious in planning review because they can make an uncertain record appear resolved. These errors are not merely technical failures; they are audit signals that show where professional verification remains necessary. For this reason, PlanTrace-VL should be understood as an evidence-grounding assistant that surfaces candidate evidence and candidate regulatory clauses, not as a substitute for professional code interpretation or public decision-making. Future validation should therefore prioritize real review packets, jurisdiction-specific rules, and planner-in-the-loop assessment, rather than treating larger models or larger benchmarks as sufficient on their own.

CONCLUSION

This study presented PlanTrace-VL, which represents planning review as an auditable chain from localized evidence and extracted facts to candidate clauses, structured notes, and human-verification status. UrbanPlanTrace-Review provides a proxy setting, and two-stage adaptation trains the model to produce this chain.

The results support a specific, bounded conclusion. PlanTrace-VL achieves 95.67% Evidence Localization and 92.89% Fact F1, showing that domain adaptation can strengthen two intermediate outputs that planners need to inspect. The same evaluation also exposes material weaknesses: Consistency F1 remains 19.66%, the unsupported-link rate is 19.67%, and insufficient visual evidence is poorly calibrated. PlanTrace-VL should therefore be interpreted as an evidence-grounding assistant, not an autonomous approval system, definitive legal compliance checker, or substitute for professional judgment.

The principal limitation is external validity: public proxy resources overrepresent technical drawings relative to municipal packets. Future work should evaluate jurisdiction-specific rules, diverse documents, cross-page aggregation, and planner-in-the-loop review on authentic case packets. Useful planning AI should make evidence easier to inspect and challenge, not conceal it behind fluent answers.

REFERENCES

- Amershi, S., Weld, D., Vorvoreanu, M., Fournery, A., Nushi, B., Collisson, P., Suh, J., Iqbal, S., Bennett, P. N., Inkpen, K., Teevan, J., Kikin-Gil, R., and Horvitz, E. (2019). Guidelines for human-ai interaction. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*.
- Bai, S., Chen, K., Liu, X., Wang, J., Ge, W., Song, S., Dang, K., Wang, P., Wang, S., Tang, J., et al. (2025a). Qwen2.5-vl technical report. *arXiv preprint arXiv:2502.13923*.
- Bai, S., Chen, K., Liu, X., Wang, J., Ge, W., Song, S., Dang, K., Wang, P., Wang, S., Tang, J., et al. (2025b). Qwen3-vl technical report. *arXiv preprint arXiv:2511.21631*.
- Blau, T., Fogel, S., Ronen, R., Golts, A., Ganz, R., Ben Avraham, E., Aberdam, A., Tsiper, S., and Litman, R. (2024). Gram: Global reasoning for multi-page vqa. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 15598–15607.
- Brinkley, C. and Stahmer, C. (2024). What is in a plan? using natural language processing to read 461 california city general plans. *Journal of Planning Education and Research*, 44(2):632–648.
- Buçinca, Z., Malaya, M. B., and Gajos, K. Z. (2021). To trust or to think: Cognitive forcing functions can reduce overreliance on ai in ai-assisted decision-making. In *Proceedings of the ACM on Human-Computer Interaction*.
- Cobbe, J., Lee, M. S. A., and Singh, J. (2021). Reviewable automated decision-making: A framework for accountable algorithmic systems. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 598–609.
- Danish, M., Munir, M. A., Shah, S. R. A., Kuckreja, K., Khan, F. S., Fraccaro, P., Lacoste, A., and Khan, S. (2025). Geobench-vlm: Benchmarking vision-language models for geospatial tasks. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 7132–7142.
- Demir Altintas, N. and Ilal, S. M. (2022). A review of automated compliance checking applications in the construction industry. *Journal of Information Technology in Construction*, 27:514–536.
- Dimyadi, J. and Amor, R. (2013). Automated building code compliance checking: Where is it at? In *Proceedings of the 19th CIB World Building Congress*.
- Fan, Z., Zhu, L., Li, H., Chen, X., Zhu, S., and Tan, P. (2021). Floorplancad: A large-scale cad drawing dataset for panoptic symbol spotting. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 10128–10137.
- Feng, J., Zhang, J., Xu, F., Zhou, C., Zhao, B., Li, Z., Chen, X., Yang, M., Liu, Y., Jin, D., and Li, Y. (2024). Citybench: Evaluating the capabilities of large language model as world model. *arXiv preprint arXiv:2406.13945*.
- Fu, X. (2024). Natural language processing in urban planning: A research agenda. *Journal of Planning Literature*.
- Fu, X., Wang, R., and Li, C. (2024). Can chatgpt evaluate plans? *Journal of the American Planning Association*, 90(3):525–536.
- Fu, X., Zhai, W., and Li, S. (2023). Using natural language processing to read plans: A study of 78 resilience plans from the united states. *Journal of the American Planning Association*, 89(1):107–119.
- Geertman, S. C. M. and Stillwell, J. (2004). Planning support systems: An inventory of current practice. *Computers, Environment and Urban Systems*, 28(4):291–310.
- Gemma Team, Kamath, A., Ferret, J., Pathak, S., Vieillard, N., Merhej, R., Perrin, S., Matejovicova, T., Ramé, A., Rivière, M., et al. (2025). Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*.
- Guan, T., Liu, F., Wu, X., Xian, R., Li, Z., Liu, X., Wang, X., Chen, L., Huang, F., Yacoob, Y., Manocha, D., and Zhou, T. (2024). Hallusionbench: An advanced diagnostic suite for entangled language hallucination and visual illusion in large vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14375–14385.
- Hettiarachchi, S. et al. (2025). Code-accord: A regulation and reasoning dataset for compliance checking. *Scientific Data*, 12:126.
- Kalervo, A., Ylioinas, J., Häikiö, M., Karhu, A., and Kannala, J. (2019). Cubicasa5k: A dataset and an improved multi-task model for floorplan image analysis. *arXiv preprint arXiv:1904.01920*.

- Klosterman, R. E. (1997). Planning support systems: A new perspective on computer-aided planning. *Journal of Planning Education and Research*, 17(1):45–54.
- Kuckreja, K., Danish, M. S., Naseer, M., Das, A., Khan, S., and Shah, F. S. (2023). Geochat: Grounded large vision-language model for remote sensing. *arXiv preprint arXiv:2311.15826*.
- Kunz, L., Klostermeier, M., Thanabalan, K., Legler, T., and Ruskowski, M. (2025). Techmb: Exploring the potential of vision language models for interpreting technical drawings. In *DS 140: Proceedings of the 36th Symposium Design for X (DFX2025)*, pages 179–188. The Design Society.
- Lai, S., Ning, Y., Yuan, Z., Chen, Z., and Liu, H. (2025). Ustbench: Evaluating language models for urban spatio-temporal prediction and reasoning. *arXiv preprint arXiv:2505.17572*.
- Le, M., Kim, R. P., Ananny, M., and Bernstein, M. S. (2024). Evaluative ai: Toward ai-generated evidence for and against decisions. *arXiv preprint arXiv:2402.01292*.
- Li, Y., Du, Y., Zhou, K., Wang, J., Zhao, X., and Wen, J.-R. (2023). Evaluating object hallucination in large vision-language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 292–305. Association for Computational Linguistics.
- Liu, F., Chen, D., Guan, Z., Zhou, X., Zhu, J., Zhou, J., and Zhou, L. (2023). Remoteclip: A vision language foundation model for remote sensing. *arXiv preprint arXiv:2306.11029*.
- Lobry, S., Marcos, D., Murray, J. D. W., and Tuia, D. (2020). Rsvqa: Visual question answering for remote sensing data. *IEEE Transactions on Geoscience and Remote Sensing*.
- Luo, R., Liu, Z., Cheng, T., Wang, J., Wang, T., Wei, X., Wang, H., Li, Y., Chai, F., Cheng, F., Ye, S., Wang, W., Zhang, Y., Qiao, Y., Zhang, H., and Zhao, X. (2025). Archcad-400k: A large-scale cad drawings dataset and new baseline for panoptic symbol spotting. In *Advances in Neural Information Processing Systems*.
- Mathew, M., Karatzas, D., and Jawahar, C. V. (2021). Docvqa: A dataset for vqa on document images. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 2200–2209.
- Miller, T. (2023). Explainable ai is dead, long live explainable ai! hypothesis-driven decision support. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency*, pages 333–342.
- Mistral AI (2025). Mistral small 3.2. Mistral AI model card, <https://docs.mistral.ai/models/model-cards/mistral-small-3-2-25-06>. Accessed 2026-06-27.
- Mistral AI Team (2025). Mistral small 3. Mistral AI news, <https://mistral.ai/news/mistral-small-3/>. Accessed 2026-06-27.
- Mleczo, J. and Desmond, M. (2023). The national zoning and land use database: A new resource for understanding urban land use regulation. *Urban Studies*, 60(14):2683–2693.
- Mosqueira-Rey, E., Hernández-Pereira, E., Alonso-Ríos, D., Bobes-Bascarán, J., and Fernández-Leal, Á. (2023). Human-in-the-loop machine learning: A state of the art. *Artificial Intelligence Review*, 56:3005–3054.
- National Zoning Atlas (2026). National zoning atlas. Accessed 2026-06-27.
- Peng, Z.-R., Lu, K.-F., Liu, Y., and Zhai, W. (2024). The pathway of urban planning ai: From planning support to plan-making. *Journal of Planning Education and Research*, 44(4):2263–2279.
- Rezvanifar, A., Cote, M., and Albu, A. B. (2020). Symbol spotting on digital architectural floor plans using a deep learning-based framework. *arXiv preprint arXiv:2006.00684*.
- Solihin, W. and Eastman, C. (2015). Classification of rules for automated bim rule checking development. *Automation in Construction*, 53:69–82.
- Talen, E. (2012). *City Rules: How Regulations Affect Urban Form*. Island Press.
- Van Landeghem, J., Tito, R., Borchmann, L., Pietruszka, M., Joziak, P., Powalski, R., Jurkiewicz, D., Coustaty, M., Anckaert, B., Valveny, E., Blaschko, M., Moens, S., and Stanislawek, T. (2023). Document understanding dataset and evaluation (dude). In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 19528–19540.
- Xu, F., Zhang, J., Gao, C., Feng, J., and Li, Y. (2023). Urban generative intelligence: A foundational platform for agents in embodied city environment. *arXiv preprint arXiv:2312.11813*.
- Yao, Y., Yu, T., Zhang, A., Wang, C., Cui, J., Zhu, H., Cai, T., Li, H., Zhao, W., He, Z., et al. (2024). Minicpm-v: A gpt-4v level mllm on your phone. *arXiv preprint arXiv:2408.01800*.
- Zhang, J. and El-Gohary, N. M. (2016). Semantic nlp-based information extraction from construction regulatory documents for automated compliance checking. *Journal of Computing in Civil Engineering*,

30(2):04015014.

- Zhang, J. and El-Gohary, N. M. (2017). Semantic-based logic representation and reasoning for automated regulatory compliance checking. *Journal of Computing in Civil Engineering*, 31(1):04016037.
- Zhang, W., Cai, M., Zhang, T., Zhuang, Y., and Mao, X. (2024). Earthgpt: A universal multimodal large language model for multisensor image comprehension in remote sensing domain. *IEEE Transactions on Geoscience and Remote Sensing*, 62:1–20.
- Zheng, Y., Liu, L., Lin, Y., Feng, J., Zhang, G., Jin, D., and Li, Y. (2025). Urbanplanbench: A comprehensive urban planning benchmark for evaluating large language models. *arXiv preprint arXiv:2504.21027*.
- Zhou, B., Yang, H., Chen, D., Ye, J., Bai, T., Yu, J., Zhang, S., Lin, D., He, C., and Li, W. (2024a). Urbench: A comprehensive benchmark for evaluating large multimodal models in multi-view urban scenarios. *arXiv preprint arXiv:2408.17267*.
- Zhou, Z., Lin, Y., Jin, D., and Li, Y. (2024b). Large language model for participatory urban planning. *arXiv preprint arXiv:2402.17161*.
- Zhu, H., Su, J., Chen, M., Wang, W., Deng, Y., Chen, G., and Zhang, W. (2025). Plangpt-vl: Enhancing urban planning with domain-specific vision-language models. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: Industry Track*, pages 2461–2483. Association for Computational Linguistics.